

Boston

Import libraries for data manipulation, machine learning, and visualization

```
import pandas as pd          # For data handling using DataFrames
import numpy as np           # For numerical operations
import tensorflow as tf      # For building and training neural networks# Import required libraries
import pandas as pd          # Pandas for data manipulation and loading CSVs
import numpy as np           # NumPy for numerical operations
import tensorflow as tf      # TensorFlow for building and training models
from sklearn.model_selection import train_test_split # Used to split data into training and testing sets
import matplotlib.pyplot as plt # Used to visualize training progress
import warnings              # Used to manage warning messages
warnings.filterwarnings('ignore') # Ignore warnings in output
```

Load the IMDB dataset from CSV

```
data = pd.read_csv('imdb_dataset.csv') # Load dataset from CSV file
print(f"Original shape: {data.shape}") # Print initial shape of the dataset
```

Preprocessing (Optional)

```
print(data.isnull().sum()) # Count and print null values per column
print(data.duplicated().sum()) # Print total number of duplicate rows
# Drop rows with missing values in 'review' or 'sentiment', and remove duplicates based on 'review'
data = data.dropna(subset=['review', 'sentiment']).drop_duplicates(subset=['review'])
print(f"Shape after removing Null & duplicates: {data.shape}") # Print shape after cleaning
```

Encode labels ('positive' -> 1, 'negative' -> 0)

```
data['sentiment'] = data['sentiment'].map({'positive': 1, 'negative': 0}) # Map string labels to integers
data['sentiment'] = (data['sentiment'] == 'positive').astype(int) # Ensures only 1s and 0s (redundant but safe)
print(data['sentiment'].value_counts()) # Print number of positive and negative reviews
```

Separate features and target

```
X = data['review'].values # Feature: reviews (text)
y = data['sentiment'].values # Target: sentiment (0 or 1)
```

Split the data into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42) # random_state ensures reproducibility
print(f"Training set shape: {X_train.shape} Test set shape: {X_test.shape}") # Print split shapes
```

Text vectorization

```
max_words = 10000 # Max number of unique words to consider
max_len = 200 # Max sequence length (padding/truncation)
vectorizer = tf.keras.layers.TextVectorization( # Create text vectorization layer
    max_tokens=max_words, # Limit vocabulary size
    output_mode='int', # Convert text to integers
    output_sequence_length=max_len # Fix sequence length to 200 tokens
)
vectorizer.adapt(X_train) # Learn the vocabulary from training text
```

Build the neural network model

```
model = tf.keras.Sequential([ # Create a sequential model
    vectorizer, # First layer: text vectorization
    tf.keras.layers.Embedding(max_words, 16, input_length=max_len), # Word embedding (dense vector of size 16)
    tf.keras.layers.GlobalAveragePooling1D(), # Pooling layer that averages over word embeddings
    tf.keras.layers.Dense(16, activation='relu'), # Hidden dense layer with ReLU activation
    tf.keras.layers.Dense(1, activation='sigmoid') # Output layer with sigmoid for binary classification
])
```

Compile the model

```
model.compile(
    optimizer='adam', # Adam optimizer for efficient gradient descent
    loss='binary_crossentropy', # Binary classification loss function
    metrics=['accuracy'] # Track accuracy during training and evaluation
)
```

Train the model

```
history = model.fit(
    X_train, y_train, # Input features and labels
    epochs=10, # Train for 10 full passes over the training data
    batch_size=32, # Train in mini-batches of 32 samples
    validation_split=0.2, # Use 20% of training data used for validation
    verbose=1 # Display training progress
)
```

Evaluate the model

```
test_loss, test_acc = model.evaluate(X_test, y_test, verbose=0) # Evaluate on test set (no output)
print(f"\nTest Loss: {test_loss:.4f}") # Print test loss (rounded to 4 decimal places)
print(f"Test Accuracy: {test_acc:.2%}") # Print test accuracy as percentage
```

Plot results

```
plt.figure(figsize=(12, 4)) # Set figure size for plots
```

Plot training and validation loss

```
plt.subplot(1, 2, 1) # First subplot
plt.plot(history.history['loss'], label='Training Loss') # Plot training loss over epochs
plt.plot(history.history['val_loss'], label='Validation Loss') # Plot validation loss over epochs
plt.title("Model Loss") # Title for loss graph
plt.xlabel('Epoch') # X-axis label
plt.ylabel('Loss') # Y-axis label
plt.legend() # Add legend to distinguish lines
```

Plot training and validation accuracy

```
plt.subplot(1, 2, 2) # Second subplot
plt.plot(history.history['accuracy'], label='Training Accuracy') # Training accuracy
plt.plot(history.history['val_accuracy'], label='Validation Accuracy') # Validation accuracy
plt.title("Model Accuracy") # Title for accuracy graph
plt.xlabel('Epoch') # X-axis label
plt.ylabel('Accuracy') # Y-axis label
plt.legend() # Add legend
```

IMDB

Import libraries for data manipulation, machine learning, and visualization

```
from sklearn.model_selection import train_test_split # To split dataset into training and testing sets
from sklearn.preprocessing import StandardScaler # To normalize features before training
import matplotlib.pyplot as plt # For plotting training metrics and predictions
import warnings # To manage warning messages
warnings.filterwarnings('ignore') # Ignore any warnings to keep output clean
# Load the Boston Housing dataset from CSV
data = pd.read_csv('Boston.csv') # Read dataset into a pandas DataFrame
data.shape # Display the number of rows and columns in the dataset
```

Preprocessing (Optional)

```
print(data.isnull().sum()) # Print number of missing values in each column
data = data.dropna() # Remove rows with any missing values

print(data.duplicated().sum()) # Print number of duplicate rows
data = data.drop_duplicates() # Remove duplicate rows from the dataset
```

```
print(f"Shape after removing Null & duplicates: {data.shape}") # Print new shape of cleaned dataset
```

```
print(data.columns) # Display column names of the dataset
# Assuming the target column is named 'medv' (adjust if different)
X = data.drop('medv', axis=1).values # Features: drop the target column 'medv'
y = data['medv'].values # Target: the column 'medv' as a NumPy array
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, # Features and target
    test_size=0.2, # 20% of data for testing
    random_state=42 # Seed to ensure reproducible split
)
```

Scale the features

```
scaler = StandardScaler() # Create a StandardScaler instance for normalization
X_train_scaled = scaler.fit_transform(X_train) # Fit to training data and scale
X_test_scaled = scaler.transform(X_test) # Scale test data using same parameters
```

Build the neural network model

```
model = tf.keras.Sequential([ # Sequential model (layers added in order)
    tf.keras.layers.Dense(64, activation='relu', input_shape=(X_train.shape[1],)), # First hidden layer with 64 neurons
    tf.keras.layers.Dense(32, activation='relu'), # Second hidden layer with 32 neurons
    tf.keras.layers.Dense(1) # Output layer for regression (1 neuron, linear activation)
])
```

Compile the model

```
model.compile(
    optimizer='adam', # Optimizer: Adam (adaptive gradient descent)
    loss='mse', # Loss function: Mean Squared Error (good for regression)
    metrics=['mae'] # Metric to track: Mean Absolute Error (easier to interpret)
)
```

Train the model

```
history = model.fit(
    X_train_scaled, y_train, # Training features and labels
    epochs=100, # Train for 100 iterations through the dataset
    batch_size=32, # Use mini-batches of 32 samples
    validation_split=0.2, # 20% of training data used for validation
    verbose=1 # Show progress bar during training
)
```

Evaluate the model

```
test_loss, test_mae = model.evaluate(X_test_scaled, y_test, verbose=0) # Evaluate on unseen test data
print(f"\nTest Mean Absolute Error: ${test_mae:.2f}k") # Print the MAE formatted in thousands
```

Make predictions

```
predictions = model.predict(X_test_scaled) # Predict house prices using the test data
```

Plot training history

```
plt.figure(figsize=(12, 4)) # Create a figure with 2 subplots side by side
```

Plot loss (MSE) over epochs

```
plt.subplot(1, 2, 1) # First subplot
plt.plot(history.history['loss'], label='Training Loss') # Plot training loss
plt.plot(history.history['val_loss'], label='Validation Loss') # Plot validation loss
plt.title("Model Loss") # Title of plot
plt.xlabel('Epoch') # Label x-axis
plt.ylabel('Loss') # Label y-axis
plt.legend() # Add legend
```

Plot MAE over epochs

```
plt.subplot(1, 2, 2) # Second subplot
plt.plot(history.history['mae'], label='Training MAE') # Plot training MAE
plt.plot(history.history['val_mae'], label='Validation MAE') # Plot validation MAE
plt.title("Model MAE") # Title of plot
plt.xlabel('Epoch') # Label x-axis
plt.ylabel('MAE') # Label y-axis
plt.legend() # Add legend
```

Plot actual vs predicted prices

```
plt.figure(figsize=(8, 6)) # New figure for actual vs predicted plot
plt.scatter(y_test, predictions, alpha=0.5) # Scatter plot of actual vs predicted prices
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--') # Diagonal reference line
plt.xlabel('Actual Price') # Label x-axis
plt.ylabel('Predicted Price') # Label y-axis
plt.title('Actual vs Predicted Prices') # Plot title
plt.tight_layout() # Automatically adjust subplot parameters
plt.show() # Display the plots
```

```
Fashion
# Import required libraries
import pandas as pd          # Used for data manipulation and analysis
import numpy as np          # Provides support for arrays and numerical operations
import tensorflow as tf     # TensorFlow library for deep learning
from tensorflow import keras # Keras API from TensorFlow for building neural networks
import matplotlib.pyplot as plt # Used for plotting graphs and visualizations
import warnings             # Used to handle warning messages
warnings.filterwarnings('ignore') # Ignore warnings in the output

# Load train and test data from CSV files
train_data = pd.read_csv('fashion-mnist_train.csv') # Load training dataset
test_data = pd.read_csv('fashion-mnist_test.csv')  # Load testing dataset
print(f"Original train shape: {train_data.shape}, test shape: {test_data.shape}") # Print shape of datasets

# Preprocessing (Optional)
print(train_data.isnull().sum()) # Print count of missing values per column in train dataset
print(test_data.isnull().sum()) # Print count of missing values per column in test dataset

train_data = train_data.dropna() # Drop rows with missing values in training dataset
test_data = test_data.dropna() # Drop rows with missing values in test dataset

print(train_data.duplicated().sum()) # Print number of duplicate rows in train dataset
print(test_data.duplicated().sum()) # Print number of duplicate rows in test dataset

train_data = train_data.drop_duplicates() # Drop duplicate rows in training dataset
test_data = test_data.drop_duplicates() # Drop duplicate rows in test dataset

print("shape after removing Null & duplicates") # Print after-cleaning info
print(f"Train : {train_data.shape}") # Print cleaned train shape
print(f"Test : {test_data.shape}") # Print cleaned test shape

# Extract features and labels
x_train = train_data.drop('label', axis=1).values # Features (images) from training data
y_train = train_data['label'].values # Labels (targets) from training data
x_test = test_data.drop('label', axis=1).values # Features (images) from test data
y_test = test_data['label'].values # Labels (targets) from test data

# Normalize pixel values
x_train = x_train.astype('float32') / 255.0 # Convert pixel values to float and normalize (0-1)
x_test = x_test.astype('float32') / 255.0 # Normalize test set similarly

# Reshape for CNN input (28x28 images with 1 channel)
x_train = x_train.reshape(-1, 28, 28, 1) # Reshape train data to 4D tensor for CNN
x_test = x_test.reshape(-1, 28, 28, 1) # Reshape test data similarly
print(f"Train shape: {x_train.shape}, Test shape: {x_test.shape}") # Print new shape

print(train_data['label'].nunique()) # Print number of unique classes (should be 10)
# Class names for visualization (Optional)
class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', # Human-readable class names
               'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle Boot']

# Visualize sample images
plt.figure(figsize=(10, 10)) # Set figure size for plotting
for i in range(25): # Plot first 25 images
    plt.subplot(5, 5, i+1) # Create 5x5 grid
    plt.xticks([]) # Remove x-axis ticks
    plt.yticks([]) # Remove y-axis ticks
    plt.grid(False) # Remove grid lines
    plt.imshow(x_train[i], cmap=plt.cm.binary) # Display image in grayscale
    plt.xlabel(class_names[y_train[i]]) # Label each image with its class name
plt.savefig('sample_images.png') # Save the plot as an image file

# Build the CNN model
model = keras.Sequential() # Create a Sequential model
    keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)), # Conv layer with 32 filters, 3x3
kernel
    keras.layers.MaxPooling2D((2,2)), # Max pooling layer with 2x2 pool size
    keras.layers.Dropout(0.25), # Dropout layer for regularization (25% dropout)
    keras.layers.Conv2D(64, (3,3), activation='relu'), # Second conv layer with 64 filters
    keras.layers.MaxPooling2D((2,2)), # Second pooling layer
    keras.layers.Dropout(0.25), # Dropout layer
    keras.layers.Conv2D(128, (3,3), activation='relu'), # Third conv layer with 128 filters
    keras.layers.Flatten(), # Flatten output into 1D vector for Dense layers
    keras.layers.Dense(128, activation='relu'), # Dense layer with 128 units and ReLU activation
    keras.layers.Dropout(0.25), # Dropout layer
    keras.layers.Dense(10, activation='softmax') # Output layer with 10 classes and softmax for
classification
)

# Compile the model
model.compile(optimizer='adam', # Adam optimizer for adaptive learning rate
              loss='sparse_categorical_crossentropy', # Loss function for multi-class classification
              metrics=['accuracy']) # Metric: classification accuracy

# Train the model
history = model.fit(x_train, y_train, # Training data
                   epochs=10, # Train for 10 epochs
                   batch_size=32, # Use batches of 32 samples
                   validation_data=(x_test, y_test), # Use test data for validation during training
                   verbose=1) # Show training progress

# Evaluate the model on test set
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=0) # Get test loss and accuracy
print(f"\nTest Loss: {test_loss:.4f}") # Print formatted test loss
print(f"Test Accuracy: {test_acc:.2%}") # Print accuracy as percentage

# Plot training and validation accuracy and loss
plt.figure(figsize=(10, 5)) # Create a new figure for plots

plt.subplot(1, 2, 1) # First subplot for accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy') # Plot training accuracy
plt.plot(history.history['val_accuracy'], label='Validation Accuracy') # Plot validation accuracy
plt.xlabel('Epochs') # Label x-axis
plt.ylabel('Accuracy') # Label y-axis
plt.title("Training and Validation Accuracy") # Title of the plot
plt.legend() # Add legend

plt.subplot(1, 2, 2) # Second subplot for loss
plt.plot(history.history['loss'], label='Training Loss') # Plot training loss
plt.plot(history.history['val_loss'], label='Validation Loss') # Plot validation loss
plt.xlabel('Epochs') # Label x-axis
plt.ylabel('Loss') # Label y-axis
plt.title("Training and Validation Loss") # Title of the plot
plt.legend() # Add legend
```